

PREMIUM RESOURCE · AI & FINANCE AUTOMATION

Finance Data Quality Checklist

A 25-item checklist for assessing and improving finance data quality before AI deployment — naming conventions, completeness, accuracy, consistency, and governance.

25 Data Quality Checks

Assessment Scoring

Common Error Patterns

Premium · €35

25

quality checks

5

categories

8

error patterns

4

priority tiers

WHY THIS MATTERS BEFORE, NOT AFTER

AI does not fix bad data — it produces confident-sounding output from it, which is a worse failure mode than a spreadsheet error, because a wrong AI-generated number often looks exactly as polished as a correct one. Run this checklist before selecting or piloting an AI tool, on the specific data source that tool will actually touch, not as a one-time exercise on the whole finance function at once.

SECTION 1

The 25-Item Checklist

Five categories of five checks each — score every item 1 (rarely true) to 5 (consistently true) as you work through your own data.

01

Naming & Structural Conventions

Whether your data is labelled and structured consistently enough for a tool to rely on it.

1.1	Consistent chart of accounts across entities The same expense type is coded identically everywhere, not three different ways across three subsidiaries.	Score (1–5)
1.2	Consistent date formats throughout One format used consistently (e.g. YYYY-MM-DD), not a silent mix of regional formats that misreads day and month.	Score (1–5)
1.3	Consistent currency and unit representation Amounts stored with a clear, consistent currency field — not implied by context or mixed within one column.	Score (1–5)
1.4	Consistent categorical labels Department, cost centre, and category codes match a single agreed list, not near-duplicates like “Marketing” and “Marketing ”.	Score (1–5)
1.5	No duplicate or conflicting identifiers Each vendor, customer, or account has exactly one ID — not two records for the same entity under slightly different names.	Score (1–5)

02

Completeness

Whether the data that should be there actually is, with nothing silently missing.

2.1	No missing required fields in transaction records Every transaction has the fields a downstream process actually depends on — not blanks quietly defaulting to zero or null.	Score (1–5)
2.2	No gaps in date sequences Every expected period is present — a missing month is a red flag, not just an inconvenience for a forecasting model.	Score (1–5)
2.3	All accounts have a corresponding description or mapping No orphaned account codes with no name, category, or owner attached.	Score (1–5)
2.4	All transactions have a corresponding audit trail record Approvals, source documents, or authorisation records exist for what governance and audit will eventually ask about.	Score (1–5)
2.5	No orphaned records An invoice without its referenced PO, or a payment without its invoice, is a gap worth finding before an AI tool tries to match them.	Score (1–5)

03

Accuracy

Whether the values themselves are actually correct, not just present and well-labelled.

3.1 Subledger reconciles to the general ledger	AR, AP, and fixed asset subledgers tie out to their GL control accounts within an explainable tolerance.	Score (1–5)
3.2 No obvious outliers or typos	A transaction 100x the typical size for its category is either genuinely unusual or a data entry error — confirmed, not assumed.	Score (1–5)
3.3 Correct sign conventions throughout	Debits, credits, and positive/negative values follow one consistent convention, not a per-analyst habit that varies file to file.	Score (1–5)
3.4 Spreadsheet formulas match their displayed values	No hardcoded numbers silently overriding a formula, invisible until someone changes an input and the total doesn't move.	Score (1–5)
3.5 Currency conversion is accurate and current	Exchange rates used are the ones actually agreed for reporting purposes, not a stale rate copied forward for convenience.	Score (1–5)

04

Consistency & Single Source of Truth

Whether the same figure means the same thing, and shows the same value, everywhere it appears.

4.1 The same figure matches across systems	Revenue in the ERP, the FP&A model, and the board deck agree — or any difference is explained, not silently tolerated.	Score (1–5)
4.2 Key metrics are defined consistently	“Revenue,” “headcount,” and other key terms mean the same thing to every team using them, documented somewhere real.	Score (1–5)
4.3 Version control exists on shared spreadsheets	One current file, clearly labelled — not five copies with names like “final_v3_ACTUAL.xlsx” circulating by email.	Score (1–5)
4.4 Period-end cutoff treatment is consistent	Every entity applies the same rule for which transactions belong to which period, not entity-specific informal habits.	Score (1–5)
4.5 Intercompany eliminations are handled consistently	The same elimination logic and timing is applied every period, not rebuilt from memory each close.	Score (1–5)

05

Accessibility & Governance

Whether the data can actually be used safely by the tools and people that need it.

5.1 Data is technically accessible to the tools that need it Not locked in a format, location, or permission structure that quietly blocks the connection an AI tool would need.	<i>Score</i> (1–5)
5.2 Clear ownership is assigned per data source Someone specific, not “finance in general,” is accountable for each source's quality.	<i>Score</i> (1–5)
5.3 Access permissions are current, not just historical Permissions reflect who should have access today, reviewed periodically — not grants nobody has revisited in years.	<i>Score</i> (1–5)
5.4 Refresh frequency matches decision-making needs Data driving a weekly decision is refreshed at least weekly — not month-old figures feeding a tool that assumes it's current.	<i>Score</i> (1–5)
5.5 A change history exists for key data sources Significant changes to a chart of accounts, mapping table, or key file are logged somewhere, not just remembered by one person.	<i>Score</i> (1–5)

SECTION 2

Assessment Scoring

How to turn 25 individual scores into a single, honest picture of readiness — by category and overall.

Scoring Method

Score each of the 25 items 1–5, giving each category a score out of 25 and an overall total out of 125. Score honestly and, where possible, with more than one person's perspective — the person closest to a data source often has the clearest view of exactly where it breaks down.

Category Scores

Category	Score (out of 25)
Naming & Structural Conventions	[score]
Completeness	[score]
Accuracy	[score]
Consistency & Single Source of Truth	[score]
Accessibility & Governance	[score]
TOTAL (out of 125)	[score]

Interpreting Your Total

Score Range	Band	What It Means
100–125	Strong	Data quality is a genuine strength — proceed with tool selection and pilot planning with confidence.
75–99	Adequate	Solid in most areas with specific, identifiable gaps — fix your lowest-scoring category before piloting on that data.
50–74	Fragile	Real risk of an AI tool amplifying existing problems — address the lowest 2 categories before proceeding.
Below 50	Not Ready	Fix foundational data issues first — an AI pilot here is very likely to produce output nobody trusts.

Read category scores, not just the total: A strong Accuracy score can sit right alongside a weak Governance score in the same total — always check which specific category is dragging the number down before deciding what to fix first.

SECTION 3

Common Error Patterns

Eight specific, recognisable patterns that account for most finance data quality problems — and why each one matters more once AI is involved.

Pattern	Why It Happens	Why It Matters for AI
Inconsistent chart of accounts across entities	Each entity's finance team built its own coding scheme independently, years apart.	Any consolidation or cross-entity AI analysis inherits the mismatch silently.
Copy-paste values instead of live formulas	A prior period's figure gets pasted as a static value "just this once" and never updated again.	The number looks current but quietly stops moving — nothing alerts anyone.
Manual "plug" entries with no clear source	A reconciliation doesn't balance, and someone posts a plug entry to make it match.	The underlying discrepancy is masked, not resolved — and an AI tool has no way to know the total is not fully explained.
Mixed date formats (DD/MM vs. MM/DD)	Files sourced from systems configured for different regional settings are combined without normalising the format.	Dates before the 13th of the month silently misread as the wrong day, with no error thrown.
Hidden characters or trailing whitespace in text fields	Data pasted from a PDF or web export carries invisible characters not visible in a normal cell view.	Exact-match lookups silently fail to match records that look identical to a human eye.
Merged cells in spreadsheets meant for export	A report was formatted for human readability first, with structured export as an afterthought.	Automated parsing breaks or misaligns columns the moment the file is opened by a tool instead of a person.
Currency values stored as text, not numbers	A field imported from a system that formats amounts with a currency symbol as part of the text string.	Sums silently return zero or an error, and few teams check every total against an independent expectation.
Duplicate vendor or customer records	The same vendor gets entered twice under slightly different names over time ("ABC Corp" vs. "ABC Corporation").	Spend analysis and matching logic both undercount or fail to connect related transactions.

Common pitfall: Assuming a data quality issue is "known" and therefore harmless because the team has quietly worked around it for years. A human workaround that corrects for a known error every month is invisible to an AI tool reading the same raw data without that same tribal knowledge.

SECTION 4

Remediation Priority Guide

Not every issue deserves the same urgency — a simple framework for deciding what to fix first, next, and later.

The Four Priority Tiers

Tier	Criteria	Approach
Fix First	High severity, low effort	Breaks something important, but is quick to correct — the best return on a small amount of effort.
Plan For It	High severity, high effort	Genuinely important to fix, but needs proper scoping — don't attempt as a quick side task.
Fix When Convenient	Low severity, low effort	Worth correcting, but not urgent — batch these in alongside other work.
Monitor, Don't Chase	Low severity, high effort	Rarely worth dedicated effort unless it starts blocking something specific — revisit if priorities change.

Assessing Severity

A data quality issue is high severity when it affects a figure used in an actual decision, a required control, or an AI use case already planned — not simply because it looks untidy. A messy but never-used column matters far less than an inconsistent chart of accounts feeding your consolidation.

- ☐ Does this issue affect a number someone currently relies on for a real decision?
- ☐ Does this issue block or undermine a specific AI use case already on your roadmap?
- ☐ Would fixing this take hours, or would it require a genuine project with its own scope?
- ☐ Is this a one-time clean-up, or a recurring problem that will simply reappear next period?

Fix recurring problems at the source, not just the symptom: A monthly manual clean-up of the same import error is a sign the fix belongs upstream — in the source system's export settings or the naming convention itself — not in an ever-repeating spreadsheet correction step.

WHERE TO START

Run this checklist specifically on the data source behind your first AI pilot, not the whole finance function at once. A focused, honest assessment on one real use case will tell you far more than a broad but shallow pass across everything.

DATA QUALITY ASSESSMENT SUPPORT

For hands-on support running a data quality assessment on your own systems before an AI rollout — a bespoke engagement is available on request.

info@mddbusinessconsulting.com