

PREMIUM RESOURCE · AI & FINANCE AUTOMATION · FLAGSHIP GUIDE

Finance AI Implementation Playbook

The complete guide to implementing AI in a finance function — readiness assessment through governance, tool selection, and a 90-day adoption plan.

Readiness Assessment Framework

Tool Selection Matrix

Governance Templates

90-Day Adoption Plan

8

parts

30

chapters

5

templates

61

pages

HOW TO USE THIS PLAYBOOK

This is a reference to work through in order the first time, then return to by chapter as your implementation progresses. Part I helps you understand where your finance function genuinely stands today; Parts II through VII take you through strategy, tool selection, governance, people, and the first 90 days; Part VIII collects every template in one place for ongoing use. Treat the frameworks as starting points to adapt — no finance function's AI implementation looks exactly like another's.

CONTENTS

Table of Contents

Thirty chapters across eight parts — use this to jump directly to what you need.

Part I — Foundations

- 1. Executive Summary
- 2. The Finance AI Landscape
- 3. Readiness Assessment Framework

Part II — Strategy

- 4. Defining Your AI Strategy
- 5. Identifying and Prioritising Use Cases
- 6. Building the Business Case

Part III — Tool Selection

- 7. The Finance AI Tool Categories
- 8. Tool Selection Matrix
- 9. Build vs. Buy vs. Extend
- 10. Vendor Evaluation and RFP Essentials

Part IV — Data & Governance

- 11. Data Readiness for AI
- 12. Governance Framework
- 13. Risk Management for AI in Finance
- 14. Regulatory and Compliance Considerations

Part V — People & Change

- 15. Change Management for AI Adoption
- 16. Roles, Responsibilities, and the AI Centre of Excellence
- 17. Training and Capability Building
- 18. Managing Resistance and Building Genuine Buy-In

Part VI — Implementation — The 90-Day Plan

- 19. Pilot Programme Design
- 20. The 90-Day Adoption Plan
- 21. Scaling From Pilot to Production
- 22. Common Implementation Pitfalls

Part VII — Measurement & Sustaining Momentum

- 23. KPIs and Success Metrics
- 24. Governance Cadence and Ongoing Oversight
- 25. Continuous Improvement and Scaling Further

Supplementary Reference

- Capstone — Putting It All Together: A Worked Case Study
- Reference — Frequently Asked Questions

Part VIII — Templates & Tools Appendix

- 26. Readiness Assessment Template
- 27. Tool Selection Matrix Template
- 28. AI Governance Charter Template
- 29. 90-Day Plan Template
- 30. Glossary of Terms

PART I

Foundations

Before choosing a tool or writing a business case, get honest about why this matters for your specific function and exactly where it stands today. This part sets the foundation everything else in this playbook builds on.

1. Executive Summary
2. The Finance AI Landscape
3. Readiness Assessment Framework

CHAPTER 1

Executive Summary

Why finance functions are adopting AI now, what actually changes, and how to use the rest of this playbook.

Finance functions are being asked to produce deeper insight, faster, with the same headcount — while the tools available to help have matured faster than most teams have had time to properly evaluate them. This playbook exists because most AI adoption in finance fails for reasons that have very little to do with the AI itself: unclear ownership, a business case built on hope rather than measurement, a rollout that skips readiness assessment entirely, or governance retrofitted after something has already gone wrong.

None of that is a reason to wait. It is a reason to be deliberate.

What Actually Changes When Finance Adopts AI Well

Before	After
Manual reconciliation and data entry consume the majority of close week.	Routine reconciliation and entry are substantially automated; close week is spent on exceptions and analysis.
Forecasts are built from static, backward-looking spreadsheets.	Forecasts update from live or near-live data, with scenario modelling built in.
Reporting requests queue up behind whoever has time to build them.	Standard reporting is self-service; the team's time goes to non-standard, judgment-heavy requests.
AI tool use, where it exists, is ad hoc and ungoverned.	AI tool use is deliberate, governed, and measured against real outcomes.

What This Playbook Assumes

This playbook assumes you are a finance leader or a member of a finance team who owns some part of this decision — not necessarily a technologist, and not necessarily starting from scratch. It assumes your organisation has at least basic Microsoft 365 or equivalent tooling in place, and it assumes you want a genuinely honest treatment of what AI can and cannot yet do reliably in a finance context, not a vendor pitch.

How to use this playbook: Read Part I in full first, even if you're tempted to skip to tool selection. An honest readiness assessment changes what “good” looks like for everything that follows — a function at Foundational maturity needs a different first 90 days than one already at Developing.

KEY TAKEAWAYS

- AI adoption in finance fails more often from unclear ownership and a missing measurement plan than from the technology itself.
- Read Part I fully before jumping to tool selection — readiness changes what “good” looks like for everything that follows.

CHAPTER 2

The Finance AI Landscape

The distinct categories of AI relevant to finance, and why conflating them is the single most common strategic mistake in this space.

“AI” is not one thing, and treating it as one thing is where many finance AI strategies go wrong before they start. The categories below solve genuinely different problems, need different governance, and often come from different vendors entirely.

Category	What It Actually Does	Finance Example
Generative AI assistants	Conversational tools that draft, summarise, and answer questions in natural language, often embedded in existing productivity tools.	Microsoft 365 Copilot drafting a variance commentary or summarising an email thread.
Workflow automation / RPA	Rule-based automation of repetitive, structured tasks — no judgment, just consistent execution of defined steps.	Power Automate routing an invoice for approval based on its amount.
Machine learning / predictive models	Statistical models trained on historical data to forecast, classify, or detect anomalies.	A cash flow forecasting model, or an anomaly detection flag on unusual transactions.
Embedded ERP / platform AI	AI features built directly into existing finance systems by the vendor, often the lowest-effort entry point.	Native anomaly flags or forecasting features inside Dynamics 365 or NetSuite.
Custom-built AI solutions	Purpose-built models or applications developed for a specific, often complex, use case no off-the-shelf tool covers well.	A custom model for matching non-standard invoice formats to purchase orders.

Why the Distinction Matters Practically

- 1

Each category carries different risk

A generative AI drafting error is embarrassing; an unreviewed predictive model driving a real financial decision is a different category of risk entirely.
- 2

Each category needs different governance

Workflow automation needs DLP-style connector controls; predictive models need model risk management practices; generative AI needs output review discipline.
- 3

Each category has a different realistic timeline

Embedded ERP AI and generative assistants can often be piloted in weeks; a custom predictive model is typically a multi-month undertaking.

Common pitfall: Building a single “AI strategy” that treats all five categories as one initiative with one timeline, one governance model, and one success metric. They need to be planned for separately, even if they're sequenced together.

A Quick Scenario for Each Category

1 Generative AI assistant, in practice

A financial analyst pastes this month's variance data into Copilot and asks for a first-draft commentary highlighting the three largest movements — then rewrites the parts that miss real business context before it goes in the board pack.

2 Workflow automation, in practice

An invoice over \$10,000 automatically routes to the submitter's manager for approval in Teams, with the amount and vendor shown directly in the approval card — no one has to manually forward an email.

3 Predictive model, in practice

A 13-week rolling cash flow model flags that week 6 is projected to dip below the minimum cash threshold, three weeks before it would have been noticed manually.

4 Embedded ERP AI, in practice

A finance team discovers their existing ERP already includes an anomaly-detection feature for unusual transactions — unused for two years because no one had turned it on.

5 Custom-built AI, in practice

A manufacturer builds a bespoke matching model because their purchase orders arrive in eleven different supplier-specific formats no off-the-shelf tool handles well.

KEY TAKEAWAYS

- “AI” spans five genuinely distinct categories, each with different risk profiles and governance needs.
- Never build one strategy that treats generative AI, workflow automation, and predictive models as a single initiative.

CHAPTER 3

Readiness Assessment Framework

Five dimensions to score honestly before choosing a single tool — the fillable version of this assessment is in Part VIII.

A readiness assessment is not a formality before the “real” work starts — it is the single best predictor of whether an AI implementation will still be running, and trusted, twelve months from now. Score each dimension honestly, ideally with more than one person's perspective, before committing to a tool or a timeline.

The Five Dimensions

D

Data Readiness

Whether your data is accessible, structured, and trustworthy enough for AI to work with reliably.

WHAT TO LOOK FOR

- Is financial data stored in structured systems (ERP, defined spreadsheets) rather than scattered across ad hoc files?
- Can the relevant data actually be accessed by the tools you're considering, technically and permission-wise?
- Is data quality (completeness, consistency, correct types) good enough that an AI tool won't simply amplify existing errors?
- Is there a single source of truth for key figures, or do different systems disagree with each other?

T

Technology Readiness

Whether your existing systems and licensing can actually support the AI tools you're evaluating.

WHAT TO LOOK FOR

- Do you have (or can you readily obtain) the licensing tier required for the AI features you want — e.g. Power Automate Premium, Copilot?
- Are your core systems (ERP, reporting tools) modern enough to integrate with AI tools without heavy custom work?
- Is there IT capacity to support a rollout, even a lightweight one, without waiting months for a project slot?
- Do you have, or can you build, basic reporting on system usage so adoption is actually measurable later?

P**Process Readiness**

Whether your processes are documented and stable enough to automate without automating chaos.

WHAT TO LOOK FOR

- Are your core processes (close, AP, AR, forecasting) documented well enough that someone new could follow them?
- Are exceptions to the standard process rare and well understood, rather than the process being mostly exceptions?
- Has the process been stable for at least a few cycles, rather than being actively redesigned?
- Is there a named process owner who could sign off on automating this specific process?

Pe**People Readiness**

Whether the team has the skills, capacity, and genuine appetite to adopt something new.

WHAT TO LOOK FOR

- Does at least one person on the team have the confidence to learn and champion a new tool?
- Does the team have realistic capacity to learn something new, or is everyone already at full stretch?
- Is there a genuine sense that AI could make the team's work better, rather than only a sense of threat?
- Has the team successfully adopted a new tool or process change in the last year or two?

G**Governance Readiness**

Whether you have, or can quickly build, the oversight structure a real rollout needs.

WHAT TO LOOK FOR

- Is there a named owner for AI tool decisions in finance, even informally?
- Do you have any existing data governance or access control practices to build on?
- Is there an existing relationship with IT/Security that could support a DLP or access review?
- Would a new AI tool be evaluated against any consistent criteria today, or purely case by case?

Scoring and Interpreting Your Result

Score each dimension's questions 1 (rarely/never true) to 5 (consistently true), giving each dimension a score out of 20 and an overall score out of 100. The fillable version of this exact assessment, with scoring formulas, is provided in Chapter 26.

Score Range	Band	What It Means
80–100	Advanced	Ready to pursue more ambitious use cases, including custom-built solutions where justified.
60–79	Developing	Solid foundations with clear gaps — typically ready for a focused pilot in 1–2 dimensions' weakest areas first.

Score Range	Band	What It Means
40–59	Emerging	Real gaps across multiple dimensions — a narrow, low-risk pilot is the right scope, not a broad rollout.
Below 40	Foundational	Address the lowest-scoring dimension directly before any AI tool selection — usually Data or Process readiness.

Read the dimensions, not just the total: A high People Readiness score can mask a genuine Data Readiness gap for a while — but AI built on unreliable data eventually produces unreliable output, however enthusiastic the team using it. Always check your lowest-scoring dimension specifically.

KEY TAKEAWAYS

- Score all five dimensions honestly, ideally with more than one person's input, before committing to a tool.
- A high People Readiness score can mask a real Data Readiness gap — always check your lowest-scoring dimension specifically.

PART II

Strategy

Readiness tells you where you stand. Strategy decides what you actually do about it — which problems to solve first, and how you'll know if it worked.

4. Defining Your AI Strategy

5. Identifying and Prioritising Use Cases

6. Building the Business Case

CHAPTER 4

Defining Your AI Strategy

What a real AI strategy has to answer, and the three failure modes that quietly derail most of them.

An AI strategy is not a list of tools you'd like to try. It is a small number of specific answers: which problems you're solving, for whom, measured how, and owned by whom. If your strategy document can't answer those four things for its top priority, it isn't a strategy yet.

Three Failure Modes to Avoid

- 1

Technology-first thinking

Starting from “we should use Copilot more” instead of “variance analysis takes too long — what would help?” The tool should be the answer to a named problem, not the starting point.
- 2

No clear owner

A strategy owned by “finance and IT together” is a strategy owned by no one. Every initiative needs one named person accountable for its outcome, even if many people contribute.
- 3

No measurement plan

If you can't say in one sentence, before you start, how you'll know this worked, you won't be able to tell afterward either — and momentum quietly stalls once nobody can point to a result.

A Simple Strategy Framework

For each candidate initiative, answer these four questions before it goes on a roadmap:

Question	Example Answer
What specific problem are we solving?	Month-end variance commentary takes 2 full days and is inconsistent in quality across analysts.
What capability does that require?	Automated variance calculation and flagging, plus AI-assisted first-draft commentary.
Which tool category fits?	Generative AI assistant (drafting) + workflow automation (calculation and flagging).
How will we measure success?	Variance commentary time drops from 2 days to half a day within two close cycles.

Align to business strategy, not just finance's own priorities: An AI initiative that saves finance five hours a week but has no connection to a broader business priority (faster decisions, better forecasting, lower cost to serve) will struggle to hold leadership's attention past the first quarter.

KEY TAKEAWAYS

- A real strategy names the problem, the capability needed, the tool category, and the success measure — for every initiative.
- The three failure modes to avoid: technology-first thinking, no clear owner, and no measurement plan.

CHAPTER 5

Identifying and Prioritising Use Cases

Where good candidate use cases actually come from, and a scoring framework for deciding what to tackle first.

Where to Find Real Candidates

- 1

Run a time-tracking exercise

Ask the team to log roughly where their hours go for two weeks. The biggest blocks of low-judgment, repetitive time are your best early candidates.
- 2

Interview for pain, not for ideas

Ask “what’s the most tedious part of your week?” rather than “what AI tools would you like?” — people are far better at naming pain than designing solutions.
- 3

Review the existing backlog

Most teams already have a running, informal list of “we really should fix this someday” — that list is a genuine head start.
- 4

Look at what’s already flagged as a control gap

Manual processes that internal audit or external audit has already commented on are strong candidates — the case for change is partly already written.

15 Candidate Use Cases, Scored

A worked example of the Impact/Effort scoring in practice — replace with your own organisation’s candidates, but use this as a starting checklist of what’s commonly worth considering.

Use Case	Category	Impact	Effort	Tier
Automate bank & cash reconciliation	Workflow automation	4	2	Quick Win
Rolling 13-week cash flow forecast	Predictive model	4	2	Quick Win
Train finance team on AI tools & data literacy	Enablement	5	2	Quick Win
Establish AI / data governance framework	Governance	4	2	Quick Win
AI-assisted month-end close checklist & workflow	Workflow automation	4	3	Major Project
Automate AP invoice processing	Workflow + predictive	5	4	Major Project
Automate AR collections & dunning	Workflow automation	4	3	Major Project
Self-service BI dashboards for stakeholders	Embedded/platform AI	4	4	Major Project
Automate variance analysis & flux reporting	Generative AI + workflow	4	3	Major Project
Migrate to cloud ERP with native AI features	Embedded/platform AI	5	5	Major Project
Financial data warehouse / single source of truth	Foundational	5	5	Major Project

Use Case	Category	Impact	Effort	Tier
Implement continuous close practices	Process + workflow	4	4	Major Project
Copilot / generative AI for reporting drafts	Generative AI	3	2	Fill-In
E-signature & digital approval workflows	Workflow automation	3	1	Fill-In
RPA for repetitive journal entries	Workflow automation	3	3	Deprioritise

Scoring: Impact vs. Effort

Score each candidate 1–5 on Impact (how much time, risk, or quality improvement it genuinely delivers) and Effort (data readiness, integration complexity, change management load). This is the same Impact/Effort logic used throughout this category's other templates — consistency here matters more than the exact scoring mechanics.

Quadrant	Definition	What to Do
Quick Win	Impact ≥ 4, Effort ≤ 2	Start here. Builds momentum and organisational trust before harder initiatives.
Major Project	Impact ≥ 4, Effort ≥ 3	Worth doing, but needs proper resourcing and its own mini business case.
Fill-In	Impact ≤ 3, Effort ≤ 2	Nice to have — slot in around higher-priority work, don't lead with these.
Deprioritise	Everything else	Low return for the effort required — revisit only if circumstances change.

Common pitfall: Letting the loudest voice in the room set priority instead of the scoring. A structured Impact/Effort exercise exists precisely to stop the most enthusiastic stakeholder's pet project from jumping the queue ahead of a genuinely higher-value quick win.

KEY TAKEAWAYS

- Find candidates through time-tracking and pain interviews, not by asking people to design AI solutions themselves.
- Score Impact and Effort consistently so priority is set by the framework, not by whoever is loudest in the room.

CHAPTER 6

Building the Business Case

The real cost categories, the real benefit categories, and what to do when a benefit is genuine but hard to quantify.

Cost Categories — The Full Picture

Category	Often Forgotten Because...
Licensing	The obvious one — rarely forgotten, but often under-scoped if usage grows past the initial tier.
Implementation / setup time	Internal hours spent configuring and testing rarely get costed the way external consulting hours do.
Training	Treated as a one-time event, when real capability-building needs ongoing time investment (see Chapter 17).
Ongoing governance and maintenance	The cost of monitoring, model retraining, and periodic review essentially never appears in the initial business case.

Benefit Categories — Beyond “Time Saved”

Category	How to Quantify It
Direct time savings	Hours per period × fully loaded hourly cost — the easiest benefit to defend, so lead with it.
Error reduction	Cost of rework, restatement risk, or audit findings avoided — harder to quantify precisely, but real.
Capacity redeployment	Hours freed up redirected to higher-value analysis — valuable, but only real if you actually redeploy the time, not just absorb it.
Risk and control improvement	Often the most defensible benefit in a regulated context, even when hardest to put a number on.

When a benefit is genuinely hard to quantify: Don't force a fake number onto a soft benefit — state it plainly as a qualitative benefit alongside the quantified ones. A business case with one honest “we believe this also improves X, though we can't yet measure it precisely” line is more credible than one where every single line has a suspiciously precise figure attached.

Common pitfall: Building the business case around benefits only, with no cost for ongoing governance. The first year's numbers look great; year two's maintenance and oversight costs then arrive as a surprise that undermines trust in the whole programme.

A Worked Example

A finance team automating bank reconciliation across three entities built the following simple first-year business case:

Line Item	Amount
Licensing (existing ERP tier, no new cost)	\$0
Implementation time (40 internal hours × \$75/hr)	\$3,000
Training (3 sessions, 15 hours total × \$75/hr)	\$1,125
Ongoing governance (2 hrs/month × 12 × \$75/hr)	\$1,800
TOTAL FIRST-YEAR COST	\$5,925
Time saved (6 hrs/week × 48 weeks × \$45/hr)	\$12,960
NET FIRST-YEAR BENEFIT	\$7,035

Note what this example deliberately includes that a rougher business case often omits: ongoing governance cost, at a realistic (not zero) hourly rate, right alongside the licensing and training lines — not just the headline time-savings figure.

KEY TAKEAWAYS

- Cost the full picture — licensing, implementation, training, and ongoing governance — not just the licence fee.
- State soft benefits honestly rather than forcing a fake precise number onto something genuinely hard to quantify.

PART III

Tool Selection

With priorities set, the next decision is which tool actually fits — and that decision needs a real framework, not just whichever vendor got the meeting first.

7. The Finance AI Tool Categories

8. Tool Selection Matrix

9. Build vs. Buy vs. Extend

10. Vendor Evaluation and RFP Essentials

CHAPTER 7

The Finance AI Tool Categories

A practical map of what’s actually available today in each category from Chapter 2, and where to start looking.

Revisiting the five categories from Chapter 2 with concrete tool examples — not a recommendation of any single vendor, but a map of where to start researching once you know which category your prioritised use case falls into.

Category	Where to Start Looking	Typical Licensing Path
Generative AI assistants	Microsoft 365 Copilot (if already on M365); vendor-native AI features in your accounting platform.	Per-user add-on licence, often on top of an existing subscription tier.
Workflow automation	Microsoft Power Automate; native workflow tools inside your ERP; dedicated RPA platforms for complex cases.	Per-flow or per-user premium tier; AI Builder credits consumed separately for document processing.
Predictive models	Purpose-built finance forecasting tools; ERP-native forecasting modules; a custom-built model as a last resort.	Subscription-based for off-the-shelf tools; project-based cost for anything custom.
Embedded ERP AI	Whatever your existing ERP vendor already ships — check what you’re already entitled to before buying anything new.	Often included in current tier, or a lower-cost upgrade than buying a separate tool.
Custom-built	Internal development team, or a specialist implementation partner for a well-scoped, narrow problem.	Project-based; ongoing maintenance cost must be planned for separately.

Check embedded ERP AI first, every time: It’s the most commonly skipped step in tool selection. Teams frequently evaluate three external vendors for a capability their existing ERP already includes, unused, in a licence tier they’re already paying for.

A Note on Licensing Tiers

Most finance teams already own more AI capability than they realise, sitting unused inside an existing licence tier. Before evaluating anything new, ask your IT team specifically: “what AI features are we already entitled to in our current Microsoft 365, ERP, and accounting platform licences?” This single question regularly surfaces a usable capability that would otherwise have triggered a months-long procurement process for something already paid for.

KEY TAKEAWAYS

- Check what your existing ERP and Microsoft 365 licence already include before evaluating anything new.
- Map your prioritised use case to a category first — that decides where to even start looking.

CHAPTER 8

Tool Selection Matrix

Seven evaluation criteria to score every serious candidate against — the fillable version is in Part VIII.

Score every candidate tool against the same seven criteria, 1–5, so comparisons are genuinely like-for-like rather than shaped by whichever vendor gave the most polished demo.

Criterion	What to Actually Check
Cost	Full cost including licensing, implementation, and ongoing governance — not just the sticker price per user.
Integration complexity	Does it connect to your existing systems natively, or does every integration require custom development?
Security & compliance	Data residency, access control granularity, and whether it meets your existing compliance obligations without exception.
Ease of use	Realistic learning curve for your actual team, not for a technically sophisticated power user.
Vendor stability	How long has this vendor supported finance use cases specifically, and what does their public roadmap suggest about longevity?
Scalability	Does the licensing and architecture hold up if usage grows five-fold, or does cost or performance break down?
Time to value	Realistic weeks to a first working result — distinct from marketing claims of “minutes to set up.”

Weight the criteria before you score, not after: Decide in advance whether security and compliance matter more than time to value for this specific use case — weighting criteria after seeing the scores just lets you justify whichever answer you already wanted.

Common pitfall: Scoring every tool against every criterion equally, regardless of the use case. Security and compliance should carry more weight for a tool touching real transactional data than for one only drafting internal commentary drafts a human will review anyway.

Worked Example — Selecting a Tool for Automated Bank Reconciliation

A finance team narrowed its options to three candidates for automating bank reconciliation: their ERP's native matching module (already licensed), a dedicated reconciliation platform, and a custom Power Automate flow with AI Builder.

Criterion	ERP-Native Module	Dedicated Platform	Custom Power Automate
Cost	5 (already licensed)	2 (new subscription)	3 (build + AI Builder credits)
Integration complexity	5 (native)	3 (needs connector setup)	2 (needs build and testing)
Security & compliance	4	4	3 (newly built, unproven)

Criterion	ERP-Native Module	Dedicated Platform	Custom Power Automate
Ease of use	3 (dated interface)	5	3
Vendor stability	5 (existing vendor)	4	n/a (internal build)
Scalability	3 (basic matching only)	5	3
Time to value	5 (already available)	3	2
TOTAL (out of 35)	30	26	16

The ERP-native module scored highest — not the most feature-rich option, but the fastest, cheapest, and lowest-risk path to a real result. The team piloted it first and kept the dedicated platform as a fallback if native matching proved too limited at scale.

KEY TAKEAWAYS

- Score every serious candidate against the same seven criteria for a genuinely like-for-like comparison.
- Weight the criteria to the actual stakes of the use case before scoring, not after seeing the results.

CHAPTER 9

Build vs. Buy vs. Extend

A decision framework for the question every tool selection eventually reaches.

Option	When It Fits	The Real Cost Most Teams Underestimate
Buy (off-the-shelf)	A well-understood, common problem that other finance teams also have — most use cases in this playbook fit here.	Configuration and integration effort, even for tools marketed as “no-code.”
Extend (embedded/platform AI)	The capability already exists inside a system you own, just unused or under-configured.	The discovery effort of finding out what you already have — often skipped entirely.
Build (custom)	A genuinely unusual problem no vendor has solved well, and the problem is important enough to justify ongoing maintenance.	Ongoing maintenance and the loss of institutional knowledge if the one person who built it ever leaves.

Default to Buy or Extend: Custom build should be the option you actively argue your way into, not the default — it is the option with the highest long-term maintenance burden and the fewest people around to support it when something breaks.

A Worked Decision

A mid-sized company needed to match unusually formatted purchase orders from a legacy supplier system to invoices. They considered three off-the-shelf matching tools first — none handled the specific format well without extensive configuration that approached custom-build cost anyway. They chose to build a narrow, purpose-built matching script, but budgeted ongoing maintenance hours explicitly into the decision from day one, rather than treating the build cost as the end of the investment.

Common pitfall: Choosing “build” because a demo of an off-the-shelf tool didn't handle an edge case perfectly, without checking whether that edge case could be solved with a modest configuration change instead of a full custom build.

KEY TAKEAWAYS

- Buy or extend should be the default; build should be the option you have to actively argue your way into.
- Budget real ongoing maintenance cost into any build decision from day one, not as a year-two surprise.

CHAPTER 10

Vendor Evaluation and RFP Essentials

The questions that separate a genuinely capable vendor from a well-rehearsed demo, and the red flags worth walking away from.

Questions Worth Asking Every Vendor

- ☐ Can you show this working on a realistic version of our own data, not a curated demo dataset?
- ☐ What does a genuinely failed or low-confidence result look like in your tool, and how is it surfaced to the user?
- ☐ What specifically is required from our IT/security team before this can go live, and how long does that typically take?
- ☐ Which of your current finance customers are a reasonable size/complexity match to us, and can we speak with one?
- ☐ What's your model or roadmap for handling a request we describe today that your product doesn't yet do?

Red Flags

- 1 No credible answer for failure cases**
A vendor who can only describe the happy path hasn't been asked hard questions by enough real customers yet.
- 2 Reluctance to test on your own data**
A tool that only performs well on a curated demo dataset is not a tool that's ready for your actual, messier reality.
- 3 Pricing that resists a straight answer**
If total cost at your actual expected usage takes three follow-up emails to pin down, expect the same opacity after you've signed.
- 4 No reference customer close to your size or sector**
Early-stage products aren't automatically bad, but you should know that's what you're evaluating, not discover it later.

KEY TAKEAWAYS

- Ask every vendor to demonstrate their tool on your own messy data, not a curated demo dataset.
- Reluctance to give a straight answer on total cost previews vendor behaviour after you've signed.

A Sample RFP Question Set, by Category

Adapt these to your specific use case — the categories matter more than the exact wording.

Category	Sample Questions
Functional fit	Can you demonstrate this working on a sample of our own data? What does a low-confidence or failed result look like to the end user?
Integration	Which of our existing systems does this connect to natively? What's required from our IT team before go-live?

Category	Sample Questions
Security & compliance	Where is our data processed and stored? What access controls and audit logging are available at our licence tier?
Support & roadmap	What's your typical response time for a production issue? What's on your roadmap for the next 12 months relevant to finance use cases?
Commercials	What is the full cost at our expected usage volume, including any per-transaction or overage charges? What does a typical contract term and exit clause look like?
References	Can we speak with a current customer of similar size and sector to us?

PART IV

Data & Governance

The foundation that determines whether an AI implementation is trustworthy, not just functional. Skipping this part to get to tool selection faster is the single most common regret finance leaders report a year later.

- 11. Data Readiness for AI
- 12. Governance Framework
- 13. Risk Management for AI in Finance
- 14. Regulatory and Compliance Considerations

CHAPTER 11

Data Readiness for AI

Why AI makes existing data problems more visible and more consequential, not less — and what to fix first.

AI does not fix bad data — it produces confident-sounding output from it, which is a genuinely worse failure mode than a spreadsheet formula error, because a wrong AI-generated number often looks exactly as polished as a correct one.

Common Finance Data Problems That Block AI Success

1 Inconsistent chart of accounts across entities

The same expense type coded three different ways across three subsidiaries makes any automated consolidation or analysis unreliable from the start.

2 Manual workarounds masking real data quality issues

If a spreadsheet has quietly corrected a system export error every month for two years, an AI tool reading the raw system data will not know to make the same correction.

3 No single source of truth

When the ERP, the FP&A model, and the board deck disagree on the same figure, an AI tool trained or pointed at any one of them inherits that disagreement.

4 Stale or infrequently refreshed data

A forecasting model fed month-old actuals will confidently produce a forecast that's already wrong the day it's generated.

A Basic Data Quality Checklist Before Any AI Pilot

- ☐ The data source for this use case has a single, agreed owner.
- ☐ Known data quality issues in this source are documented, not just informally known by one person.
- ☐ The refresh frequency of this data matches the decision frequency the AI tool is meant to support.
- ☐ Access permissions on this data reflect who should currently have access, not historical grants nobody has reviewed.

Common pitfall: Choosing the pilot use case with the most exciting AI capability rather than the cleanest available data. A modest use case on excellent data will outperform an ambitious one on poor data, and will do more to build organisational trust in AI generally.

KEY TAKEAWAYS

- AI doesn't fix bad data — it produces confident-sounding output from it, which is a worse failure mode than a visible error.
- Pick the pilot with the cleanest available data, not the most exciting AI capability.

CHAPTER 12

Governance Framework

The core structure every AI tool in finance needs — ownership, access control, and decision rights — regardless of which category it falls into.

Governance is not a compliance afterthought bolted on before an audit — it is what makes the difference between an AI tool the organisation trusts enough to rely on, and one that quietly gets worked around after the first visible mistake.

The Core Governance Principle

AI tools generally do not grant new access to data — they operate within whatever permission structure already exists. The real governance risk is that AI makes existing over-permissioning immediately visible and actionable, often for the first time, by summarising or surfacing content a user technically had access to but never had reason to actually see.

Governance Structure — Who Owns What

Decision	Typical Owner
Which AI tools are approved for use in finance	Finance leadership, in coordination with IT/Security.
Data access and permission reviews	IT/Security, with finance data owners flagging what needs review.
Model or output accuracy standards for a specific use case	The finance process owner for that specific area (e.g. the controller for close-related tools).
Ongoing usage monitoring and incident response	A named individual or small governance group — see Chapter 16 for the Centre of Excellence model.

Governance Checklist Before Broad Rollout

- ☐ A Data Loss Prevention (DLP) or equivalent policy governs which systems an automation tool can connect across.
- ☐ New tools are piloted in a non-production environment or with a small group before being rolled out broadly.
- ☐ Every AI tool or automation in active use has a named, current owner.
- ☐ A quarterly review checks for orphaned tools, stale permissions, and unreviewed usage patterns.

Governance is easier to build early than to retrofit: A governance structure designed before the first tool goes live takes a fraction of the effort of retrofitting one after three ungoverned pilots are already running — and retrofitting always happens under worse conditions, usually right after something has gone wrong.

KEY TAKEAWAYS

- AI tools don't grant new access — they make existing over-permissioning immediately visible and actionable.
- Build governance before the first tool goes live; retrofitting it always happens under worse conditions.

CHAPTER 13

Risk Management for AI in Finance

Adapting model risk management — a discipline banking has used for decades — to the AI tools finance teams are adopting now.

Financial services has managed “model risk” — the risk that a model used for a real decision is wrong, misused, or poorly understood — for decades, well before “AI” was the term in use. The core discipline transfers directly to today’s AI tools, even lightweight ones.

Risk Type	What It Actually Looks Like in Finance	A Reasonable Control
Accuracy risk	A forecasting model or AI-drafted analysis is confidently wrong, and nobody catches it before it reaches a decision-maker.	Independent review of output against a known-good baseline before relying on it for a real decision.
Over-reliance risk	A team stops applying its own judgment because “the tool said so,” including on outputs the tool was never validated for.	Explicit documentation of what a tool is, and is not, validated to be used for.
Bias risk	A model trained on historical data quietly perpetuates a past pattern (e.g. which vendors got flagged for review) that was never actually correct.	
Explainability risk	Nobody can explain why a specific transaction was flagged or a specific number was produced, when a manager or auditor asks.	Logging the inputs and logic behind a decision separately from the tool’s own run history (see Chapter 14).

A Lightweight Risk Register — Worked Example

You do not need a full financial-services model risk function for a departmental AI tool. You do need, at minimum, a simple log like the one below. Three tools with a one-page risk note each is far better governance than a single detailed template nobody actually fills in.

Tool	Owner	What Could Go Wrong	What Would Catch It
Invoice AI extraction	AP Manager	Misreads an amount on a poor-quality scan, posting the wrong figure.	Amounts over \$5,000 route to manual review before posting.
Cash flow forecast model	Treasury Lead	Model trained on pre-disruption data understates a real risk period.	Monthly variance check against actuals; model flagged for review if variance exceeds 15%.
Variance commentary AI draft	FP&A Manager	Plausible-sounding but factually wrong explanation goes into a board pack unreviewed.	Every AI-drafted commentary requires named human sign-off before distribution.
PO-matching exception bot	Controller	Legitimate exceptions get lost in a high volume of low-priority alerts.	Weekly review of exception log completion rate, not just alert volume.

Common pitfall: Applying rigorous model risk management to a low-stakes drafting assistant while leaving a high-stakes predictive model with no review process at all, because the drafting tool happened to roll out first and got the attention. Match the rigor of the control to the actual stakes of the decision.

KEY TAKEAWAYS

- Match the rigor of your control to the actual stakes of the decision the AI tool supports.
- A one-page risk note per tool, actually filled in, beats one exhaustive template nobody completes.

CHAPTER 14

Regulatory and Compliance Considerations

What changes, from a controls and audit perspective, the moment a previously manual process becomes AI-assisted.

SOX and Internal Control Implications

Automating a control does not remove it from scope — it changes what the control evidence looks like. A manual approval leaves an email trail; an automated approval flow needs its own evidence that the logic is correct and hasn't silently changed. Treat automating an existing SOX-relevant control as a control change requiring its own review, not simply a productivity improvement.

Data Privacy

Any AI tool touching personal data (employee expense data, customer payment information) inherits your existing data privacy obligations. Confirm data residency and processing location match your requirements before assuming a tool is compliant simply because it's widely used elsewhere in the organisation.

Audit Trail Requirements

A tool's own run history or activity log typically confirms that something executed — it does not, on its own, explain why a specific decision was made in terms an auditor can independently verify. Where a genuine audit trail matters, log the decision and its inputs to a separate, durable record deliberately, rather than relying on a tool's built-in logging alone.

Involve internal audit early, not at year-end: A five-minute conversation with internal audit before a control is automated is far cheaper than a finding raised after the fact, once the old manual evidence trail no longer exists and the new one wasn't designed with audit needs in mind.

A Pre-Automation Compliance Checklist

- ☐ Internal audit has been made aware this control is being automated, before it happens, not after.
- ☐ The new evidence trail has been reviewed against what the old manual evidence trail actually provided.
- ☐ Data residency and processing location have been confirmed to meet existing privacy obligations.
- ☐ A clear answer exists for “how would we explain this specific decision to an external auditor” before go-live.

KEY TAKEAWAYS

- Automating a control changes what the evidence looks like — it doesn't remove the control from scope.
- A tool's run history confirms execution, not necessarily an auditable reason for a specific decision.

PART V

People & Change

Every framework so far assumes people will actually adopt what you build. This part is about making that true — the part of implementation most technically-focused plans skip.

15. Change Management for AI Adoption

16. Roles, Responsibilities, and the AI Centre of Excellence

17. Training and Capability Building

18. Managing Resistance and Building Genuine Buy-In

CHAPTER 15

Change Management for AI Adoption

ADKAR, adapted specifically for finance — because adoption succeeds or fails at the level of daily habits, not announcements.

Awareness, Desire, Knowledge, Ability, and Reinforcement — ADKAR — is the individual change model this playbook uses throughout, because AI adoption in finance is ultimately a series of individual decisions to use a new tool instead of a familiar habit, repeated daily under real deadline pressure.

A Awareness

IN ONE LINE

Each person can explain, specifically, why this tool now — not a vague sense of “modernising.”

D Desire

IN ONE LINE

Each person can answer “what’s in this for me,” not just for the organisation.

K Knowledge

IN ONE LINE

People know how to use this specific tool for their specific tasks, not AI in the abstract.

A Ability

IN ONE LINE

People can apply it under real deadline pressure, not just in a clean training sandbox.

R Reinforcement

IN ONE LINE

The new behaviour still sticks three months after launch week’s energy has faded.

Go deeper when you need to: This chapter is a summary — the full ADKAR treatment for finance, with communication templates, a training design guide, an adoption measurement toolkit, and a resistance management protocol, is covered in this category's dedicated AI Change Management Guide for Finance Teams.

KEY TAKEAWAYS

- ADKAR failure almost always concentrates in Desire or Reinforcement, not in Awareness.
- The full change management toolkit for finance — templates, training design, and resistance protocol — lives in the dedicated Change Management Guide.

CHAPTER 16

Roles, Responsibilities, and the AI Centre of Excellence

Who actually does what, and when a small, informal Centre of Excellence earns its keep.

Core Roles

Role	Responsibility
Executive sponsor	Owns the overall AI direction for finance and removes organisational blockers — does not need to be technical.
Initiative owner	One named person per tool or use case, accountable for its outcome and its governance.
Power users / champions	Early adopters who help train peers and surface real-world edge cases the initial design missed.
IT / Security partner	Owns technical feasibility, access control, and DLP-style policy enforcement.
Process owners	Confirm a specific automated process still reflects how the business actually needs it to work.

When to Formalise a Centre of Excellence

A small finance team running one or two pilots doesn't need a formal Centre of Excellence — clear individual ownership is enough. Once three or more tools are live across different processes, a lightweight CoE (even a monthly one-hour meeting among initiative owners) earns its keep by catching duplicated effort, inconsistent governance, and tools quietly drifting out of alignment with each other.

Common pitfall: Standing up a formal Centre of Excellence with a charter, a budget, and quarterly reporting before a single pilot has actually gone live. Governance structure should grow with real usage, not precede it by a year.

A Sample RACI for a Single Initiative

Activity	Responsible	Accountable	Consulted	Informed
Tool selection	Initiative owner	Executive sponsor	IT/Security	Wider finance team
Access & DLP configuration	IT/Security	IT/Security lead	Initiative owner	Initiative owner
Training delivery	Power users	Initiative owner	Process owner	Pilot team
Go/no-go scaling decision	Initiative owner	Executive sponsor	CoE (if formed)	Wider finance team

KEY TAKEAWAYS

- Every initiative needs one named, accountable owner — shared ownership functions as no ownership at all.
- Formalise a Centre of Excellence once three or more tools are genuinely live, not before the first one launches.

CHAPTER 17

Training and Capability Building

A four-stage training structure built around real workflows, not generic feature tours.

- 1

Concept session

30–45 minutes on what the tool does, why now, and what won't change — no hands-on practice yet.
- 2

Hands-on sandbox session

60–90 minutes practising on a copy of real, non-critical data, building a shared internal reference as you go.
- 3

Real-task paired practice

The first 2–4 real uses, paired with a confident peer, sequencing lower-stakes tasks before higher-stakes ones.
- 4

Office hours and ongoing support

A standing, predictable time slot for questions — not just “ask if you need help,” which almost nobody takes up voluntarily.

Segment by persona, not just by role: Early adopters, skeptics, and less digitally confident team members each need a different version of the same training — skeptics specifically should be invited to scrutinise, not managed around, since their concerns often catch real issues.

A Sample Two-Week Training Schedule

Day	Session
Day 1	Concept session (all pilot team members, 45 minutes).
Day 3	Hands-on sandbox session, split into two smaller groups of 3–4 people.
Day 5–10	Paired real-task practice, scheduled individually around each person's actual workload.
Day 12	First office hours session — same time slot every week going forward.

KEY TAKEAWAYS

- Train on real workflows and real data, segmented by role — never a generic feature tour.
- Ongoing office hours matter more to lasting adoption than the polish of the original training session.

CHAPTER 18

Managing Resistance and Building Genuine Buy-In

Why resistance in finance is often more reasonable than it first looks, and how to respond to it well.

Common, Reasonable Sources of Resistance

- ☐ Job security fear — especially acute given AI's visible impact on routine transactional work.
- ☐ Professional scepticism toward unverified numbers — a trained strength in finance, not a flaw to train away.
- ☐ Fear of being blamed for an AI tool's mistake once they've signed off on its output.
- ☐ A working process feeling safer than a faster, unfamiliar one under real deadline pressure.

The Response Protocol

Identify whether resistance is a skill gap, a fear-based concern, or a legitimate risk concern — each needs a genuinely different response. A skill gap needs more or different training. A fear-based concern needs to be addressed directly, not argued away with reassurances the person doesn't believe. A legitimate risk concern needs to be taken seriously — it may reveal a real control gap leadership hadn't fully considered.

Common pitfall: Treating all resistance as a communication problem to be overcome. Some resistance is protecting something worth protecting — a control point, a review step. The honest answer is sometimes that the resistance is right, and the process needs to change, not the person.

Matching the Response to the Type

Type	What It Sounds Like	The Right Response
Skill gap	"I don't really understand how to use this yet."	More or different training — paired practice, not a repeat of the same session.
Fear-based	"I'm worried this means my role won't exist in a year."	Address directly and honestly — don't repeat a reassurance the person doesn't believe.
Legitimate risk concern	"What happens if it gets this wrong and I've already signed off on it?"	Take it seriously — it may reveal a real control gap worth fixing, not just calming.

KEY TAKEAWAYS

- Diagnose whether resistance is a skill gap, a fear, or a legitimate risk concern before responding to it.
- Some resistance is correct — the honest answer is sometimes that the process needs to change, not the person.

PART VI

Implementation — The 90-Day Plan

Strategy, tools, governance, and people are all in place on paper. This part is where it actually happens — a real, week-by-week plan for the first 90 days.

19. Pilot Programme Design

20. The 90-Day Adoption Plan

21. Scaling From Pilot to Production

22. Common Implementation Pitfalls

CHAPTER 19

Pilot Programme Design

Choosing a pilot narrow enough to control risk, but meaningful enough to actually prove something.

A pilot that's too broad becomes a full rollout with none of the governance a full rollout needs. A pilot that's too narrow proves nothing anyone will trust as evidence for scaling. The right pilot sits deliberately in between.

1 Choose one process, one team, one clear success measure

Not “improve finance efficiency” — something as specific as “reduce bank reconciliation time for the EMEA entity from 3 days to 1.”

2 Define success criteria before you start, in writing

Written down before launch, success criteria can't quietly shift after the fact to match whatever the result turns out to be.

3 Pick a pilot team that includes at least one skeptic

A pilot team of only enthusiasts will surface fewer real problems than one that includes someone who will genuinely push back.

4 Timebox it explicitly

90 days, reviewed formally at the end — not an open-ended “we'll see how it goes” that never actually reaches a decision point.

Use your readiness assessment and priority scoring together: The best pilot candidate is usually a Quick Win (Chapter 5) in a process where your Chapter 3 readiness assessment showed genuine strength — not necessarily the single highest-impact idea on the list.

A One-Paragraph Pilot Charter

“We will pilot [tool] for [specific process] with [named team], for 90 days, with success defined as [specific, measurable outcome]. [Named executive] is the sponsor, [named person] is the initiative owner, and at Day 90 we will formally decide to scale, adjust, or stop based on the criteria above.”

Write this down and share it before Week 1 starts. A pilot without this sentence written down somewhere is far more likely to drift, quietly, into an ungoverned permanent fixture nobody formally decided to keep.

KEY TAKEAWAYS

- Write success criteria down before you start, so they can't quietly shift after the fact to match the outcome.
- Include at least one genuine skeptic on the pilot team — their scrutiny catches real problems.

CHAPTER 20

The 90-Day Adoption Plan

A detailed, week-by-week plan across three phases — the fillable version is in Part VIII.

This plan assumes a single, well-scoped pilot (Chapter 19) and a readiness assessment already completed (Chapter 3). Adjust timing to your own context — the sequencing matters more than the exact week numbers.

Phase 1 — Foundation (Days 1–30)

Week	Focus	Key Activities
Week 1	Confirm scope and sponsorship	Finalise pilot use case and written success criteria; secure an executive sponsor; confirm the initiative owner.
Week 2	Score and select the tool	Complete the Tool Selection Matrix (Chapter 8) for this specific use case; confirm licensing and access requirements with IT.
Week 3	Set up governance basics	Confirm DLP/access policy covers this tool; set up a non-production test environment; complete the data quality checklist (Chapter 11).
Week 4	Prepare the team	Finalise the pilot team; schedule training; communicate the “why now” to the wider team (Awareness stage, Chapter 15).

Watch for in Phase 1: Vague or missing written success criteria at the end of Week 1 is the single strongest early predictor that a pilot will drift. Don't let Week 2 start without them genuinely locked down.

Phase 2 — Build & Test (Days 31–60)

Week	Focus	Key Activities
Week 5	Concept training and initial build	Run the concept training session (Chapter 17); begin configuring the tool against the pilot process.
Week 6	Hands-on training and sandbox testing	Run the sandbox session on non-critical historical data; refine configuration based on what breaks.
Week 7	First real-task use	Begin paired real-task practice on the lowest-stakes version of the process; gather first-round feedback.
Week 8	Refine and expand	Address feedback from Week 7; expand real-task use to slightly higher-stakes instances; start tracking usage metrics.

Watch for in Phase 2: A quiet productivity dip during Weeks 5–7 is normal, not a sign of failure — the team is learning a new tool on top of, not instead of, their existing workload for now. Expect it and plan around it rather than treating it as an early red flag.

Phase 3 — Stabilise & Decide (Days 61–90)

Week	Focus	Key Activities
Week 9	Full pilot-team usage	The full pilot team uses the tool for real work; run the first pulse-check survey (adoption sentiment); address emerging resistance directly.
Week 10	Mid-point review	Check progress against the written success criteria; adjust governance or training if gaps have appeared.
Week 11	Steady-state operation	Office hours running on a predictable schedule; usage and outcome data reviewed against baseline.
Week 12	Formal 90-day review	Score the pilot against its original success criteria; decide to scale, adjust and extend, or stop; document what was actually learned.

Watch for in Phase 3: Usage that quietly declines after Week 9's initial full rollout, even while sentiment survey scores stay positive, is the earliest warning sign of a stalled Reinforcement stage (Chapter 15) — check actual usage data, not just how people say they feel about the tool.

The Week 12 decision must include “stop” as a genuine option: A pilot that hasn't met its written success criteria by Week 12 should be allowed to end without it being treated as a personal or team failure — that honesty is what makes the next pilot's success criteria credible.

KEY TAKEAWAYS

- A quiet productivity dip during Weeks 5–7 is normal, not an early warning sign of failure.
- “Stop” must be a genuine option at the Week 12 review, not just scale or adjust.

CHAPTER 21

Scaling From Pilot to Production

What genuinely changes when a successful pilot becomes a permanent, wider-reach tool.

What Changes at Scale

1 Governance rigor increases

A tool ten people use occasionally needs lighter oversight than the same tool as the default process for a whole department.

2 Edge cases multiply

More users and more real-world variation surface exceptions the pilot's narrower scope never encountered.

3 Training needs a repeatable structure

Ad hoc, informal training that worked for a five-person pilot team doesn't scale to onboarding new hires eighteen months from now.

Go/No-Go Checklist Before Scaling

- ☐ The pilot met its written success criteria, not a retrospectively adjusted version of them.
- ☐ Governance controls (Chapter 12) have been reviewed for the larger scale, not just carried over unchanged from the pilot.
- ☐ A repeatable training path exists for onboarding someone who wasn't part of the original pilot team.
- ☐ A named owner is confirmed for the tool at its new, larger scope — not necessarily the same person as the pilot owner.

Scale in stages, not all at once: Moving from a 5-person pilot directly to a 200-person rollout skips the chance to catch scale-specific problems (heavier IT load, more edge cases, more varied skill levels) while the group affected is still small enough to recover from a mistake quickly.

KEY TAKEAWAYS

- Governance rigor, edge cases, and training needs all increase at scale — none of it simply carries over unchanged.
- Scale in stages rather than jumping straight from a small pilot to a full rollout.

CHAPTER 22

Common Implementation Pitfalls

A consolidated retrospective of where finance AI implementations most often go wrong — across every part of this playbook.

1 Skipping the readiness assessment to move faster

The single most common root cause of a stalled implementation, discovered only after the fact, six months in.

2 A business case built entirely on time savings

With no accounting for governance and maintenance cost, meaning year-two costs arrive as an unpleasant surprise.

3 Choosing the tool before naming the problem

Technology-first thinking that produces an impressive pilot with no one quite sure what business problem it solved.

4 Governance retrofitted after a visible mistake

Rather than designed in from the start, when it would have taken a fraction of the effort.

5 Training treated as a single one-time event

Rather than an ongoing capability-building process with real support after the session ends.

6 No written success criteria before the pilot starts

Which quietly lets “success” be redefined after the fact to match whatever result actually occurred.

Revisit this chapter at your Week 12 review: Most of these pitfalls are only obvious in hindsight — checking this list explicitly during your 90-day review (Chapter 20) is more reliable than hoping you'd notice one happening to you in real time.

KEY TAKEAWAYS

- Revisit this pitfall list explicitly at your Week 12 review — most of these are only obvious in hindsight.
- The most common root cause across nearly all of them is skipping Part I's readiness assessment to move faster.

PART VII

Measurement & Sustaining Momentum

A successful 90-day pilot proves a concept. What happens in month four through month twelve determines whether it becomes a genuine capability or a tool that quietly stops being used.

23. KPIs and Success Metrics

24. Governance Cadence and Ongoing Oversight

25. Continuous Improvement and Scaling Further

CHAPTER 23

KPIs and Success Metrics

Three categories of metric every AI initiative in finance should track — usage, outcome, and sentiment.

Category	Example Metrics	Why It Matters
Usage	% of team actively using the tool weekly; tasks completed with the tool vs. without.	Distinguishes genuine adoption from a one-time training attendance nobody acted on.
Outcome	Time saved on the target process; error rate before vs. after; hours redeployed to higher-value work.	The actual business case for the initiative — measure it, don't just assume the projected benefit materialised.
Sentiment	Confidence self-rating; would-recommend score; open-ended concern themes from pulse surveys.	A leading indicator of whether adoption is heading toward Reinforcement (Chapter 15) or quietly stalling.

Set the baseline before you start, not after: You cannot credibly claim a process improved if you didn't measure its state before the tool existed. Baseline measurement belongs in Phase 1 of the 90-day plan (Chapter 20), not as an afterthought at the Week 12 review.

A Sample KPI Set for a Close-Automation Initiative

KPI	Baseline	Target
% of close tasks automated	20%	70%
Days to close	10	5
Manual journal entries per month	150	40
Finance hours on manual data entry / week	25	8

These four alone will not tell the whole story — pair them with at least one sentiment metric (Chapter 15) so a technically-improving KPI trend doesn't mask a team quietly disengaging from the tool that produced it.

KEY TAKEAWAYS

- Track usage, outcome, and sentiment together — a good outcome trend can mask quiet disengagement underneath it.
- Set your baseline before the tool exists, not after, or you can't credibly claim the improvement was real.

CHAPTER 24

Governance Cadence and Ongoing Oversight

What a realistic, sustainable review rhythm looks like once initial rollout excitement has faded.

Cadence	What Gets Reviewed
Monthly	Usage metrics for each active tool; any flagged incidents or unusual patterns; outstanding training needs.
Quarterly	Full governance checklist (Chapter 12) re-run; access and permission review; tool ownership confirmed still current.
Annually	Full readiness reassessment (Chapter 3) — maturity should be moving up, and this confirms whether it actually is.

Common pitfall: Running a thorough governance review at launch, then letting the cadence quietly lapse once no one is actively watching. Usage patterns, ownership, and risk all drift over time — a tool that was well-governed at launch is not automatically well-governed a year later.

A Sample Monthly Review Agenda

- ☐ Usage summary for each active AI tool since the last review — trending up, flat, or declining?
- ☐ Any incidents, near-misses, or flagged accuracy concerns raised since the last review.
- ☐ Outstanding training requests or onboarding needs for new team members.
- ☐ Any tool whose named owner has changed role — confirm reassignment has actually happened, not just been noted.

Keep the review short and recurring, not long and rare: A focused 30-minute monthly check-in that actually happens every month catches far more real problems than an exhaustive quarterly review that gets pushed back twice and finally happens five months late.

KEY TAKEAWAYS

- A short review that actually happens every month beats a thorough one that keeps getting postponed.
- Usage, ownership, and risk all drift over time, even after a genuinely well-governed launch.

CHAPTER 25

Continuous Improvement and Scaling Further

Building a genuine pipeline of future initiatives once the first one has proven itself.

1 Return to your Chapter 5 prioritisation list

The next Quick Win on that list is usually a better next step than an ambitious new idea generated by the first pilot's success.

2 Ask your pilot's power users what they'd automate next

The people closest to the tool daily usually have the sharpest instinct for what the next real opportunity actually is.

3 Re-run the readiness assessment before the next pilot

Your Data, Process, and Governance readiness will have genuinely changed since Chapter 3 — don't assume last time's scores still hold.

4 Keep the Centre of Excellence cadence (Chapter 16) light

As more initiatives launch, resist letting oversight become heavier than the actual risk of what's being governed justifies.

WHERE THIS PLAYBOOK ENDS AND YOUR WORK CONTINUES

A finance function's AI maturity is never “finished” — it's a capability that compounds, pilot by pilot, governed consistently, measured honestly. The templates in Part VIII are built to be reused for every initiative that follows this one, not just the first.

KEY TAKEAWAYS

- Return to your existing prioritisation list for the next initiative rather than chasing a new exciting idea.
- Keep Centre of Excellence oversight proportional to actual risk as more initiatives launch.

CAPSTONE

Putting It All Together — A Worked Case Study

A fictional mid-sized finance function's path through every part of this playbook, start to finish.

Meridian Manufacturing is a fictional composite based on patterns common across mid-sized finance functions — a 12-person finance team supporting a \$180M manufacturing business, still closing the books largely by hand.

Part I — Readiness Assessment (Chapter 3)

Meridian's controller ran the full readiness assessment with her team. The result: 58/100, "Emerging." Data Readiness scored lowest (9/20) — three different spreadsheet versions of the same cost centre mapping were in active use. Governance scored second-lowest (11/20) — no one owned AI tool decisions, formally or informally.

Part II — Strategy and Prioritisation (Chapters 4–6)

Rather than starting with the most exciting idea (a predictive demand-forecasting model the CFO had seen at a conference), the team scored 15 candidate use cases using the Chapter 5 framework. Bank reconciliation automation scored as a clear Quick Win: high impact, low effort, and — critically — sitting on genuinely clean data despite the function's overall Emerging maturity.

Part III — Tool Selection (Chapters 7–10)

Following Chapter 7's advice, IT confirmed Meridian's existing ERP already included a native bank-matching module, unused. It scored highest on the Tool Selection Matrix against two external alternatives — not the most feature-rich, but by far the fastest and lowest-risk path to a real result.

Part IV — Governance (Chapters 11–14)

Before piloting, the controller became the named owner for finance AI decisions — an informal designation, but a real one. A simple DLP check confirmed the matching module wouldn't expose data outside the existing ERP boundary. No new access was granted; existing permissions were reviewed and one stale admin account was removed in the process.

Part V — People (Chapters 15–18)

One team member was openly skeptical, worried the tool was a first step toward reducing headcount. The controller answered directly: no headcount decision had been made, and she said so plainly rather than offering a reassurance she couldn't back up. That skeptic became one of the pilot's most useful voices, catching two configuration issues in testing that the enthusiastic early adopters had missed.

Part VI — The 90-Day Pilot (Chapters 19–22)

Written success criteria: reduce EMEA bank reconciliation time from 3 days to 1 day within two close cycles. By Week 12, actual reconciliation time had dropped to 1.5 days — short of target, but a clear enough improvement, with a specific known cause (one bank feed still required manual CSV upload), to justify continuing rather than stopping.

Part VII — Measurement and What Came Next

Twelve months later, Meridian had automated three more processes using the same framework, each pilot faster than the last because the governance structure and training pattern no longer had to be built from scratch. The originally-exciting demand-forecasting model finally got built in month nine — by then, on far more trustworthy underlying data than it would have had as the very first initiative.

The lesson underneath the story: Meridian's success came from sequencing — an honest readiness assessment, a deliberately modest first pilot, and governance built before it was tested by scale — not from choosing a more sophisticated tool than any other finance team could access.

REFERENCE

Frequently Asked Questions

Direct answers to the questions finance leaders raise most often when starting this work.

Do we need a formal AI strategy document before doing anything?

No — you need the four answers in Chapter 4 (problem, capability, tool category, success measure) for your first pilot. A polished strategy document can follow once you have a real result to point to, not before.

What if our data quality is genuinely poor across the board?

Choose a pilot process specifically where data quality is a relative strength, even if the overall function scores Foundational or Emerging on the Chapter 3 assessment. A modest win on clean data builds more trust than an ambitious attempt on poor data.

Do we need a dedicated AI governance team?

Not at the pilot stage — one named owner is enough. A lightweight Centre of Excellence (Chapter 16) earns its place once three or more tools are genuinely live, not before.

How do we handle a team member who refuses to engage with the tool at all?

Diagnose first (Chapter 18): is this a skill gap, a fear, or a legitimate concern? Each needs a different response, and treating outright refusal as pure stubbornness without checking which one it is usually misses the real issue.

What's a realistic timeline before we see genuine results?

The 90-day plan in Chapter 20 is built to reach a real, measurable result by Day 90 for a well-scoped pilot — not full organisational transformation, which is a multi-year, multi-pilot process, not a single 90-day sprint.

Should we wait for a more mature or specialised tool before starting?

Usually no. Most finance teams already have unused AI capability in an existing licence (Chapter 7). Starting with what you already have, on a well-chosen pilot, teaches you more than waiting for a theoretically better tool that hasn't been evaluated against your actual data yet.

What happens if our first pilot fails to meet its success criteria?

Treat it as real information, not a personal failure (Chapter 20). Document what was learned, and check the Chapter 22 pitfalls list — most missed pilots trace back to one or two identifiable, fixable causes, not a fundamental flaw in the approach.

How do we know if we're over-relying on an AI tool's output?

A practical test: if no one on the team could confidently explain why the tool produced a specific result when asked, that's a sign of over-reliance risk (Chapter 13) worth addressing before it becomes a real problem.

PART VIII

Templates & Tools Appendix

Every framework from this playbook, collected here as ready-to-use templates — for this initiative and every one that follows it.

26. Readiness Assessment Template

27. Tool Selection Matrix Template

28. AI Governance Charter Template

29. 90-Day Plan Template

30. Glossary of Terms

CHAPTER 26

Readiness Assessment Template

The complete fillable version of the Chapter 3 framework — all 20 questions, score each 1–5.

Dimension	Statement	Score (1–5)
Data	Financial data is stored in structured systems, not scattered ad hoc files.	[score]
Data	The relevant data can actually be accessed by the tools being considered.	[score]
Data	Data quality is good enough that AI won't simply amplify existing errors.	[score]
Data	There is a single source of truth for key figures across systems.	[score]
Technology	We have, or can readily obtain, the licensing tier required.	[score]
Technology	Our core systems are modern enough to integrate without heavy custom work.	[score]
Technology	IT capacity exists to support a rollout without a months-long project queue.	[score]
Technology	Basic usage reporting exists, or could be built, to measure adoption later.	[score]
Process	Core processes are documented well enough for someone new to follow.	[score]
Process	Exceptions to the standard process are rare and well understood.	[score]
Process	The process has been stable for several cycles, not actively being redesigned.	[score]
Process	There is a named process owner who could sign off on automating this.	[score]
People	At least one person has the confidence to learn and champion a new tool.	[score]
People	The team has realistic capacity to learn something new right now.	[score]
People	There is genuine appetite, not only fear, about what AI could do here.	[score]
People	The team has successfully adopted a new tool or process change recently.	[score]
Governance	There is a named owner for AI tool decisions in finance.	[score]
Governance	Existing data governance or access control practices exist to build on.	[score]
Governance	There is an existing relationship with IT/Security to support an access review.	[score]
Governance	A new AI tool would be evaluated against consistent criteria today.	[score]

Scoring Summary

Dimension	Score (sum of 4 questions, out of 20)
Data Readiness	[sum]
Technology Readiness	[sum]
Process Readiness	[sum]

Dimension	Score (sum of 4 questions, out of 20)
People Readiness	[sum]
Governance Readiness	[sum]
TOTAL (out of 100)	[sum]

Use the full framework for a real assessment: This template is a quick-reference starting point. The complete framework, with all 20 questions across all 5 dimensions and full scoring bands, is in Chapter 3.

CHAPTER 27

Tool Selection Matrix Template

Score every serious candidate against the same seven criteria from Chapter 8.

Criterion	Tool A	Tool B	Tool C
Cost (1–5)	[score]	[score]	[score]
Integration complexity (1–5)	[score]	[score]	[score]
Security & compliance (1–5)	[score]	[score]	[score]
Ease of use (1–5)	[score]	[score]	[score]
Vendor stability (1–5)	[score]	[score]	[score]
Scalability (1–5)	[score]	[score]	[score]
Time to value (1–5)	[score]	[score]	[score]
TOTAL	[sum]	[sum]	[sum]

Weight before you score: Revisit Chapter 8’s guidance on weighting criteria by the actual stakes of the use case before filling this in — an unweighted total treats a drafting assistant and a transaction-level automation as equally risk-sensitive, which they are not.

CHAPTER 28

AI Governance Charter Template

A one-page starting charter — adapt the clauses below to your own organisation.

Purpose

This charter governs the evaluation, approval, and ongoing oversight of AI tools used within the finance function, including generative AI assistants, workflow automation, predictive models, and any embedded or custom-built AI capability.

Scope

Applies to any tool that processes finance data, automates a finance process, or is used by a member of the finance team in a professional capacity — including tools already in informal or trial use.

Decision Rights

Decision	Owner
Approval of a new AI tool for finance use	[named role, e.g. Finance Director]
Data access and DLP policy configuration	[named role, e.g. IT Security Lead]
Accuracy standards for a specific use case	[named role, e.g. relevant process owner]
Ongoing usage monitoring	[named role or Centre of Excellence]

Review Cadence

Monthly usage review; quarterly full governance checklist; annual readiness reassessment — see Chapter 24 for what each review covers.

CHAPTER 29

90-Day Plan Template

The fillable version of the Chapter 20 plan — add your own owner and status per week.

Week	Focus	Owner	Status
1	Confirm scope and sponsorship	[name]	[not started / in progress / done]
2	Score and select the tool	[name]	[status]
3	Set up governance basics	[name]	[status]
4	Prepare the team	[name]	[status]
5–6	Training and initial build/test	[name]	[status]
7–8	First real-task use and refinement	[name]	[status]
9–10	Full pilot usage and mid-point review	[name]	[status]
11–12	Steady-state operation and formal review	[name]	[status]

CHAPTER 30

Glossary of Terms

Plain-language definitions for every term used throughout this playbook.

Term	Definition
ADKAR	A five-stage individual change model — Awareness, Desire, Knowledge, Ability, Reinforcement — used throughout this playbook for change management.
AI Builder	Microsoft's low-code AI capability embedded in Power Automate, commonly used for document data extraction.
Anomaly detection	A predictive-model technique that flags transactions or values statistically unusual relative to historical patterns.
Centre of Excellence (CoE)	A lightweight, often informal group coordinating governance and knowledge-sharing once multiple AI initiatives are running.
Champion / power user	An early adopter who helps train peers and surfaces real-world edge cases a rollout's initial design missed.
Connector	A pre-built integration allowing a workflow automation tool to communicate with a specific system, such as SharePoint or Outlook.
DLP (Data Loss Prevention)	A policy classifying which systems and connectors can exchange data with each other, preventing unintended data flows.
Embedded AI	AI features built directly into an existing system by its vendor, rather than a separately purchased tool.
Explainability	The degree to which a specific AI-driven decision or output can be traced back to a clear, human-understandable reason.
Generative AI	AI that produces new content — text, drafts, summaries — in response to a natural language prompt.
Governance	The ownership, access control, and review structure that keeps an AI tool trustworthy over time, not just functional at launch.
Impact/Effort matrix	A prioritisation framework scoring candidate initiatives on their value (Impact) against their difficulty (Effort).
Model risk	The risk that a model used for a real decision is wrong, misused, or poorly understood — a discipline long established in financial services.
Over-reliance risk	The risk that a team stops applying its own judgment because an AI tool's output is trusted by default, including for uses it was never validated for.
Pilot	A deliberately scoped, timeboxed trial of an AI tool on a single process with a small team, before any decision to scale.
Readiness assessment	A structured scoring exercise across Data, Technology, Process, People, and Governance dimensions, done before committing to a tool.

Term	Definition
Reinforcement	The ADKAR stage concerned with a new behaviour still sticking after the initial rollout's energy has faded.
RFP (Request for Proposal)	A structured document used to gather comparable information from multiple vendors before selecting a tool.
RPA (Robotic Process Automation)	Rule-based automation of repetitive, structured tasks, with no judgment applied — workflow automation, in this playbook's terms.
SOX control	An internal control relevant to Sarbanes-Oxley compliance; automating one changes its evidence trail and requires its own review.
Success criteria	The specific, measurable outcome a pilot is judged against, written down and agreed before the pilot begins.
Use case	A specific, named problem an AI tool is meant to solve — the starting point for any tool selection, never the tool itself.
Workflow automation	Rule-based automation of a defined business process, typically built with a tool such as Power Automate.
Workflow Definition Language (WDL)	The underlying schema Power Automate and Azure Logic Apps use to define a flow's triggers and actions.

BONUS TEMPLATE

Sample Board / Leadership Update

A one-page format for reporting AI initiative progress upward, reusable across every initiative in this playbook.

Format

Section	Content
Initiative & owner	[Name of initiative, named owner]
Status	[On track / At risk / Delayed — one word, stated plainly]
Progress against success criteria	[The specific written criteria from Chapter 19/20, and current measured progress against them]
What's working	[1–2 specific, genuine points — not generic praise]
What's not working yet	[Stated honestly — this builds more credibility than an all-good report]
Decision or support needed	[Anything specifically requiring leadership input or removal of a blocker]
Next milestone	[Specific date and what will be true by then]

Keep it to one page: A board or leadership update that can't fit on one page is a sign the initiative's scope, or the report's honesty about what actually matters, needs tightening — not that the template needs to grow.

WHERE TO GO FROM HERE

Start with Chapter 3's readiness assessment this week, even informally. Everything else in this playbook — the tool you select, the pilot you design, the governance you build — will fit your actual situation better once you're honest about where your finance function genuinely stands today.

Premium resource from MDD Business Consulting. Licensed for individual professional use. Not for redistribution or resale.

FINANCE AI IMPLEMENTATION SUPPORT

For hands-on support running a readiness assessment, building a business case, or designing your own 90-day pilot — a bespoke engagement is available on request.

info@mddbusinessconsulting.com